


Unanimity flips sign under omission



Why both results are the same fact

the blind spot is a SHARED PRIOR, not uncertainty

systematic → invisible from inside

sampling spread 1.08 of 6 · logprob -0.005 · self-report 0.9
every output-distribution signal sat at chance



systematic → ensembling does not average it away

cross-model disagreement as an ambiguity signal: AUROC 0.472
more models does not mean more spread when the bias is shared



systematic → but visible from OUTSIDE, as agreement

independently trained families converging under omission means they are following the prior, not the requirement. Agreement is the alarm.